

Knodle: Modular Weakly Supervised Learning with PyTorch

Anonymous RepL4NLP 2021 submission

Abstract

Strategies for improving the training and prediction quality of weakly supervised machine learning models vary in how much they are tailored to a specific task or integrated with a specific model architecture. In this work, we propose a software framework *knodle* that treats weak data annotations, deep learning models, and methods for improving weakly supervised training as separate, modular components. The standardized interfaces between these independent parts account for data- and model-agnostic weak supervision method development, but still allow the training process to access fine-grained information such as data set characteristics, matches of heuristic rules, as well as elements of the deep learning model ultimately used for prediction. Hence, our framework can encompass a wide range of training methods for improving weak supervision, ranging from methods that only look at correlations of rules and output classes, to those that harness the interplay of neural networks and weakly labeled data. We illustrate the benchmarking potential of the framework with a performance comparison of several reference implementations on a selection of datasets that are already available in *knodle*.

1 Introduction

Most of today's machine learning success stories are built on top of huge labeled data sets. Creating and maintaining such data sources manually is a time-consuming and complicated, and thus expensive and error-prone, process. Various research directions aim to reduce the hunger for bigger and better datasets.

One of the most popular approaches that has recently gained traction is *weak supervision*. The learning algorithm is confronted with labels which are easy to obtain but lack any correctness guarantee, and as such often demand denoising. Among others, the creation of such labels is performed with regular expressions, keyword lists or external databases (the

latter is a so-called *distant supervision*). Typically, weak supervision methods and implementations are tailored towards a domain-specific task or integrated with a specific model architecture. For example, (Lin et al., 2016) introduce a special architecture called *attention over instances* for the relation extraction task, (Keith et al., 2017) propose an EM-based algorithm for event extraction and (Hedderich et al., 2021) uses the perturbation for named entity recognition. Such diversity and specificity of approaches makes it difficult to compare or transfer them across tasks without introducing any task- or data-dictated adjustments.

We introduce *knodle*: a framework for **Knowledge-supervised Deep Learning**, i.e., weak supervision with neural networks. The framework provides a simple tensor-driven abstraction based on PyTorch allowing researchers to efficiently develop their methods for improving weakly supervised machine learning models and try them interchangeably to find the one that works best for the given data. Within this work, we refer to a denoising method as any method that helps to improve the weakly supervised data regardless the type of noise and the exact denoising object (weak labels, weak rules etc).

The following points summarize *knodle*'s main design goals:

- **Data abstraction.** A tensor-driven data abstraction subsumes a large number of input variants and is applicable to diverse tasks.
- **Method independence.** A decoupled implementation of weak supervision denoising methods and prediction model enables comparability and accounts for domain-specific inductive biases.
- **Accessibility.** A high-level interface makes it easy to test existing methods, incorporate new ones and benchmark them against each other.

Several denoising algorithms are already included to *knOdle* in order to allow the “out-of-the-box” framework testing. We also discuss their performance on some datasets, also included to the *knOdle* ecosystem, each of which exemplifies different possible scenarios in weakly supervised learning, such as the amount of used labeling functions or their precision-recall balance. Moreover, depending on the properties of heuristic labelings, methods for improving weak labels need to remove spurious matches in some cases, or generalize from them in others.

It is clear that such a diverse problem space should be paired with a rich pool of methods so that the most appropriate denoising method can be found for any task or dataset. *knOdle* allows to easily explore the spaces of weakly supervised learning settings and label improvement algorithms, and hopefully will facilitate a better understanding of the phenomena that are inherent to weakly supervised learning.

The code is publicly available at [...].

2 Related work

Many strategies have been introduced to reduce the need for huge amounts of data annotated manually by experts. Among these are *active learning* (Sun and Grishman, 2012), where automatically selected instances are manually annotated by experts, and *semi-supervised learning*, (Kozareva et al., 2008; Agichtein and Gravano, 2000), where a small labeled dataset is combined with a huge unlabeled dataset. Since the arrival of BERT, pre-training and fine-tuning schemes showed good results for a moderate to small amount of labeled data. However, *knOdle* completely abstains from labeled data during training by making use of weakly supervised learning.

2.1 Weak supervision

In weak supervision, tedious expert work is replaced with easy to obtain, but potentially error-prone labels, that are usually derived from a set of heuristic rules. One of the most popular strategies of weakly supervised learning is *distant supervision*, which uses knowledge from existing data sources to annotate unlabeled data. The technique is very popular in relation extraction (Craven and Kumlien, 1999), where various knowledge databases, such as WordNet (Snow et al., 2004), Wikipedia (Wu and Weld, 2007) and Freebase (Mintz et al., 2009), are used as annotation sources.

When using heuristic rules, it is not uncommon that one sample turns out to be annotated by multiple rules. The most straightforward approach to resolve

such cases is majority voting, which is used in early weak supervision algorithms (Thomas et al., 2011) as well as in more recent experiments (Krasakis et al., 2019; Boland and Krüger, 2019). However, majority voting does not deal with the different kinds of noise introduced by weak supervision, and more noise-specific algorithms are necessary. For example, the noise produced by *incomplete labels*, which stems from the incompleteness of weak supervision sources, is commonly reduced by data manipulations, e.g. by enhancing the knowledge base (Xu et al., 2013), by a thorough construction of negative examples to balance the positive ones (Riedel et al., 2013), or by explicitly modelling missing knowledge base information with latent variables (Ritter et al., 2013). The problem of *noisy features* (increased amount of false positive labels because of overgeneralizing information from databases) is often approached by using a relaxed distant supervision assumption (Riedel et al., 2010; Hoffmann et al., 2011), by active learning with additional manual expertise (Sterckx et al., 2014), with help of topic models (Yao et al., 2011; Roth and Klakow, 2013), as well as by using a combination of multiple methods (Roth, 2014).

Apart from that, methods treat the identified potentially noisy samples differently. They are either kept for further training with reduced weights (Jat et al., 2018; He et al., 2020), corrected (Shang, 2019) or eliminated (Qin et al., 2018). Thus, denoising methods vary significantly depending on the data and task, what makes the creation of some platform, where they could be compared to each other, crucial.

2.2 Structure Learning

Structure learning assumes multiple weak labels per instance where each label is created by a so called *labeling function*. The goal is to learn a dependency structure within these labeling functions which motivates the term structure learning. Most labeling functions are generated by human intuitions, motivating correlation and dependence between labeling functions. The first algorithm was implemented in software package Snorkel (Ratner et al., 2017), which also comprised the data programming paradigm allowing to programmatically create labeling functions. It was followed by various improvements (Bach et al., 2017; Varma et al., 2019) and variations (Chatterjee et al., 2019; Maheshwari et al., 2020).

2.3 Noise-aware learning

A common idea to mitigate single noisy labels is to build an architecture which accounts for noisy

data. There are different approaches that model noise-robustness by adapting a loss function, such as a noise adaption after the loss calculation (Patrini et al., 2017), or cross-entropy and the mean absolute error generalization (Zhang and Sabuncu, 2018). Alternatively, a special noise layer in a neural network could be used (Sukhbaatar et al., 2015). Many approaches are based on noise assumptions, such as on the assumption of symmetric label noise (van Rooyen et al., 2015). Another approach aims at finding and removing wrongly labeled samples from the training procedure. Recently, the confidence learning framework CleanLab was introduced (Northcutt et al., 2021), which is based on the intuition that low-confidence predictions are more likely to be labeled wrongly. Note that most of these methods were built with the assumption that there is one label corresponding to each instance, while *knOdle* makes use several weak signals per instance.

2.4 Crowdsourcing annotations

Another possible solution to reduce the cost of manual data supervision is **crowdsourcing**. In order to increase the supervision accuracy for a task, most crowdsourcing experiments rely on annotations from multiple people, and the final label is defined by majority voting (Kosinski et al., 2012) or measuring the inter-annotator agreement (Tratz and Hovy, 2010). More sophisticated denoising strategies include anomaly detection (Eskin, 2000), annotator’s reliability modelling (Dawid and Skene, 1979), Bayesian approaches (Raykar and Yu, 2012) and generative models (Hovy et al., 2013). Some mistakes can be also corrected: for example, the ones that are consistently made by careful but biased people (Ipeirotis et al., 2010), while others, for example in case of *spammers*, are to be eliminated (Raykar and Yu, 2012).

As they both represent a task with multiple noisy labels, automatically and human labeled data resemble each other to an extent that makes their denoising algorithms possibly interchangeable, despite different annotation origins. For example, (Rehbein and Ruppenhofer, 2017) adapted the MACE algorithm (Hovy et al., 2013), initially proposed for improving noisy annotations from human annotators, to the setting of denoising automatically labeled data for named entity recognition. In a similar way, we introduce *WSCrossWeigh* (see Section 4 for more details) in order to demonstrate the usefulness of the *Knodle* framework in the experimentation in transferring the denoising algorithms from one domain to another.

2.5 Frameworks

knOdle bases upon the ideas of several software frameworks. On a low level, *knOdle* is built on top of PyTorch (Paszke et al., 2017). As for design decisions, we followed several other high-level libraries that aim to ease the training and prediction experience. Namely, we drew inspiration from PyTorch lightning (Falcon, 2019), which in essence tries to remove the burdens of writing your own train loop, and Huggingface’s Transformers library (Wolf et al., 2020), which gives easy access to various transformer-based architectures in a fixed manner, so that they can be effortlessly interchanged in code.

3 Weakly supervised learning with *knOdle*

The *knOdle* architecture provides a layer of abstraction that allows weakly supervised learning signals to be integrated into PyTorch alongside algorithms for improving them. On the one hand, since *knOdle* has access to the information about rules that matched for each sample, it is not restricted to methods that denoise only weak rules, as Snorkel (Ratner et al., 2017)), or only weak labels, as Cleanlab (Northcutt et al., 2021), - it provides a functionality to implement a denoising method that uses any of these aspects. On the other hand, a tight connection to the deep learning model enables the integration of denoising methods that use or manipulate the prediction model itself.

To the best of our knowledge, *knOdle* is the first framework offering such wide functionality. For that reason we believe that it can become a testbed where different algorithms for improving the weakly supervised data are implemented and compared with each other to find the most fruitful task-to-denoising-method combination or to use it as a foundation for further studies.

The framework follows two main design principles, outlined below:

3.1 Tensor-based representations of input data and weak label matches

Similar to Pytorch models, where the data is already assumed to be in Tensor form, and the specific pre-processing that led to a particular input representation is outside the scope of the core deep learning model, we choose to exclude the process of weak label generation from *knOdle*. Rather, we encode the weak labels in two tensors. One tensor contains information about which rules matched for each data instance, while another tensor describes the relationship between rules and output classes.

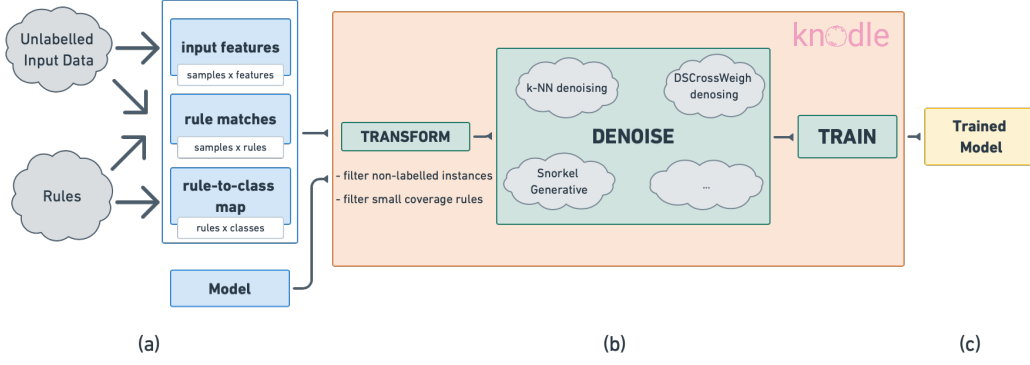


Figure 1: The figure gives an overview of our system. (a) represents the preprocessed input, given as tensors. (b) resembles the internals of *knOdle*. The *Trainer* classes introduced in section 3.4 handle transformation, denoising and model training in an integrated manner. (c) shows the output, a trainer PyTorch model.

Formally, assume we have n samples, r rules and k classes. Rule matches are gathered in a binary matrix $Z \in \{0,1\}^{n \times r}$, where $Z_{ij} = 1$ if rule j matches sample i . The initial mapping from rules to the corresponding classes is given by another binary matrix $T \in \{0,1\}^{r \times k}$, $T_{jk} = 1$ if rule j is indicative of class k .

This separation between one tensor that contains rule matches and another tensor that translates them to labels allows *knOdle* to access this fine-grained information during training for certain denoising algorithms. This is in contrast to other approaches that treat weak supervision as learning from a noisy heuristic label matrix $Y_{heur} = ZT$ without direct access to the individual rules.

3.2 Separation of the prediction model from the weak supervision aspects.

Although *knOdle* contains both methods to improve the data, such as rules denoising and generalization, and the training of standard PyTorch models, these aspects are separated in the framework. Therefore the same prediction model definition can be used for direct or weakly supervised training, allowing researchers to easily create benchmarks with consistent base models. However, even though the prediction model is defined separately, the method dealing with weak data has access to it during training. Thus, a denoising method and model training can reinforce each other in an integrated manner, as, e.g., in WSCross-Weigh method that will be covered in the next section. This is in contrast to approaches that modularize denoising and training by first adjusting label confidences with the help of rules correlations only and then training a model with the adjusted labels (Takamatsu et al., 2012; Ratner et al., 2017). Furthermore, *knOdle*'s design is much more flexible compared to

approaches where denoising is so tightly integrated into the underlying prediction model architecture that it could not be changed (Sukhbaatar et al., 2015).

3.3 Handling of negative instances

Different tasks need a different logic to handle data samples where no rule matched. These samples are traditionally called *negative instances*. For example, in binary sentiment analysis, there most likely will be labeling functions for both, the positive and the negative class, covering all the possible outcomes. In such a setting, there is no use for negative, or unlabeled, examples and filtering them out is reasonable. Meanwhile, other tasks may include an unspecific "other" class comprising the instances from none of the distinct classes, like "no relation" examples in relation extraction contained in the test set. Current weak supervision frameworks provide only one of the two options: negative samples are either filtered out (Ratner et al., 2017) or included to the training dataset (Shu et al., 2020).

knOdle includes a configurable functionality for both ways of handling such cases depending on the task at hand. From a technical point of view, there is a `filter_non_labelled` flag in a configuration file, which could be set to `False` if the samples with no rule matched samples are to be filtered out. Otherwise, to make up for the missing explicit annotation for the negative samples from the labeling functions, an additional `other_class` is defined for the negative instances, that would be treated in classifier training in the same way as other classes. This encapsulates and rounds up the weak supervision component, meaning that the denoising methods can deal with both strategies in the same way after this point, regardless of how the "other" class is handled. The exact `other_class_id` could be

either provided by user or calculated automatically by *knOdle*. The amount of negative instances that is to be included to the classifier training set can be defined specifically in each denoising algorithm.

3.4 Implementation Details

As the most popular deep learning frameworks, such as TensorFlow (Abadi et al., 2015) and PyTorch (Paszke et al., 2017), we realise learning as a mapping from input tensor(s) to output tensor(s) guided by a loss function that measures the quality of the learned mapping. However, while the most common solution is to represent the training data by a *design matrix* $X \in \mathbb{R}^{n \times d}$ (n instances represented by d feature dimensions) and a *label matrix* $Y \in \mathbb{R}^{n \times k}$ (k classes), input of *knOdle* are matrices X , Z and T described above. The heuristic labels are calculated later during the weakly supervised learning using the information contained there.

To ensure seamless use of the developed algorithms, the weakly supervised algorithms need to be tightly integrated with automatic differentiation and optimization supported by PyTorch. All of this is done within *Trainer* classes. During initialization, they receive data, a model, and a method-specific configuration, inheriting from *Config* and containing such information as model training parameters, criterion, validation method, class weights and others.

While in the standard deep learning frameworks training can be executed by calling `model.train(X, Y)`, in *knOdle* the same functionality would be invoked with the following command (illustrates the *Trainer* with k -NN search, which we describe in 4; the full code snippet is in Appendix ??):

```
KnnDenoisingTrainer(model, X, Z, T,
                    config).train()
```

4 Denoising algorithms (Trainers)

We already provide some out-of-the-box trainers which can serve as baselines or starting points for further investigations. All *Trainer* are compatible with any Pytorch model; *knOdle* provides logistic regression and HuggingFace’s *transformers* (Wolf et al., 2020) classification models as an example.

Majority Voting Baseline. As a simple baseline, the rules are directly applied to the test data. If several rules match, the prediction is done based on the majority, ties are broken randomly.

Trainer without Denoising. The simplest trained model is the *NoDenoisingTrainer*. The majority vote is computed on the training data and used to train the given model. This is the most direct use of the rule matches for training a classifier. To cover cases where several rules match, this trainer can be configured to either use a one-hot encoding of the winning label from the majority vote or a distribution over labels (relative to the number of matching rules).

Trainer with kNN Denoising. This *Trainer* includes the label denoising method with a simple geometric interpretation (*KNNDenoisingTrainer*). The intuition behind this method is that similar samples should be activated by the same rules which is allowed by a smoothness assumption on the target space. The trainer looks at the k most similar samples sorted by, for example, TF-IDF features combined with L_2 distance, and activates the rules matching the neighbors to create a denoised \hat{Z} . Importantly, *knOdle* allows separate features for the model training and the neighborhood activation. This method also provides a way to activate rules for initially unmatched samples.

Trainer with Snorkel Denoising. *knOdle* provides a wrapper of the Snorkel system (Ratner et al., 2017) *SnorkelTrainer* which incorporates both generative and discriminative Snorkel steps. The generative step constitutes a denoising method in *knOdle*’s terminology, while the discriminative step corresponds to a prediction model. The structure within labels and rules, in our notation $P(Y, Z, T)$, is learned in an unsupervised fashion by the generative model. Afterwards, the final discriminative model, i.e. the prediction model, is trained with weak labels provided by the generative model, following the general *knOdle*’s design. Also both steps are conveniently provided in a single method call. The consecutive application of k -NN and Snorkel denoising methods is also available as *SnorkelKNNDenoisingTrainer*.

Trainer with Weak Supervision CrossWeigh Denoising. Finally, we implemented our own algorithm for noise correction in weakly supervised data. It is based on the *CrossWeigh* method (Wang et al., 2019) and included to *knOdle* as *WSCrossWeighTrainer*. While the original *CrossWeigh* method was proposed for mistakes identification in crowdworkers annotations, we claim that it could be applied for denoising the weakly supervised data as well. In *WSCrossWeigh* we adopted the same logic for estimating the reliability of weakly annotated data, but made some necessarily corrections conditioned by the weak supervision learning specificity.

dataset	classes	train / test samples	rules	avg. rule hits	class ratio
Spam	2	1586 / 250	10	1.63	0.47
Spouse	2	22254 / 2701	9	0.34	0.08
IMDb	2	40000 / 5000	6786	33.97	0.50
TAC-based RE	41	1937211 / 18660	182292	0.51	-

Table 1: Summary of data statistics. The average rule hits are computed on the train set. Class ratio describes the amount of positive samples in the test set for binary classification datasets, i.e. data skewedness.

The main intuition behind WSCrossWeigh is the following: if a decision rule corresponds to a wrong class and, therefore, annotates all samples in the training set with a wrong label, the model is likely to learn the incorrect pattern and to make similar mistakes while labeling the test samples. However, if we take a sufficiently big portion of data *without* samples where this particular rule matched, train the model on it, and then classify the test samples, the predictions are supposed to contradict the initial *wrong* labels, what helps us to trace the misclassified samples and reduce their importance in final classifier training.

As in the original CrossWeigh, the basic idea is similar to the k -fold cross-validation, where input data is split into k folds, each of which becomes, in turn, a test set, while the model is trained on the others. In WSCrossWeigh, however, the splitting is performed not randomly, but basing on rules. Firstly, the rules are randomly split into K folds $\{r_1, \dots, r_k\}$ and, iteratively, each $fold_l$ is chosen to form a test set that is built from all samples matched this fold’s rules. Other samples build a training set that is used for classification model training. During the testing of the trained model on the hold-out fold samples, the predicted label \hat{y}_i for each test sample x_i is compared to the label y_i originally assigned to x_i by weak supervision. If $\hat{y}_i \neq y_i$, the sample x_i is claimed to be potentially mislabeled, and its weights w_{x_i} is reduced by a value of an empirically estimated parameter ϵ . This procedure is repeated several times with different splittings to detect misclassified samples more accurately.

The final classifier is trained on the whole reweighed training dataset. As a result, the more times the original y_i label of data sample x_i was suspected to be wrong, the smaller is its weight w_{x_i} , and, therefore, the smaller part it will play in the classifier training.

Along with other denoising algorithms, WSCrossWeigh was tested on the datasets described in Section 5 and showed quite promising results: it outperforms all other algorithms on three out of four datasets. For more details please see Section 6.

5 Datasets

Apart from denoising methods, we also include some well-known datasets in the *knOdle*-specific tensor format in order to demonstrate the abilities of our framework. All datasets, collected for straightforward prediction tasks (sentiment analysis, spam detection, and relation extraction) and being quite simple, still have their own peculiarities in Z and T matrices, that are worth investigating.

The *knOdle* ecosystem, therefore, provides modular access to datasets and denoising methods (that can, in turn, be combined with arbitrary deep learning models), enabling easy experimentation. The overview of dataset statistics is shown in Table 1.

Spam Dataset. The first dataset is a Spam dataset based on the YouTube comments dataset (Alberto et al., 2015). Here, the task is to classify whether a text is relevant to the video or holds spam, such as advertisement. The dataset has a small size of both train and test sets. Thus, a single wrongly labeled instance might have quite a big impact on the learning algorithm. We use the preprocessed version by the Snorkel team (Snorkel, 2020b). Among others, the rules were created based on keywords and regex expressions.

Spouse Dataset. This relation extraction dataset is based on the Signal Media One-Million News Articles Dataset¹ (Corney et al., 2016). The task is to decide whether a sentence holds a spouse relation or not. We also inherited labeling functions written by the Snorkel team that include a set of known spouse from DBPedia (Lehmann et al., 2014) as well as keywords and encoded language patterns. The difficulty of the Spouse dataset is its skewness: over 90% of samples in the test set hold a no-spouse relation. Again, the preprocessed version by the Snorkel team is used (Snorkel, 2020a), so the results can be related to previous studies (Ratner et al., 2017).

IMDb Dataset. The third dataset is based on the well-known IMDb dataset, which consists of short movie reviews. The task is to determine whether a

¹<http://research.signalmedia.co/newsir16/signal-dataset.html>

Mode	Spam	Spouse			IMDb	TAC-based RE		
	Acc	P	R	F1	Acc	P	R	F1
Majority vote	0.81	0.12	0.79	0.22	0.65	0.09	0.001	0.001
Majority + DistilBert	0.87	0.09	0.90	0.17	0.67	0.20	0.19	0.19
k -NN + DistilBert	0.94	0.12	0.86	0.21	0.50	0.10 [†]	0.11 [†]	0.10 [†]
WSCrossWeigh + DistilBert	0.94	0.09	0.69	0.16	0.73	0.25	0.27	0.26
Snorkel + DistilBert	0.88	0.13	0.70	0.23	-	-	-	-

Table 2: Results of the classifier training with different denoising methods on the datasets (on test sets) added to *Knodle*.
[†]The neighbors were searched with Approximate Nearest Neighbors (Bernhardsson, 2015) because of computation complexity of k -NN search.

review holds a positive or negative sentiment. Despite the training set has labels, we do not use them in our experiments, but handle this data as unsupervised. To create the Z and T matrices, we use the positive and negative keyword lists (Hu and Liu, 2004), with a total of 6800 keywords.

TAC-based Relation Extraction Dataset. Lastly, in order to evaluate the relation extraction task on a larger dataset, we added to *knodle* a dataset built over Knowledge Base Population challenges in the Text Analysis Conference. For development and test purposes the TACRED corpus annotated via crowdsourcing and human labeling from KBP (Zhang et al., 2017) is used. As human labels are not allowed in weak training, the training is performed not on the TACRED dataset, but on a weakly-supervised noisy corpus built on TAC KBP corpora (Surdeanu, 2013; Roth, 2014), which was annotated with entity pairs extracted from Freebase (Google, 2014) with corresponding relations mapped to the 41 TAC relations. The amount of entity pairs per relation was limited to 10,000 and each entity pair is allowed to be mentioned in no more than 500 sentences. An important difference of this dataset to the other three is the presence of negative instances added to the dataset in equal proportion to the positive ones.

6 Experiments

Our main aim was not to receive the best possible results for each dataset, but rather to demonstrate that the different methods for improving the weak annotations perform with a different success rate on different datasets depending on the data properties and peculiarities. That also serves to demonstrate the *knodle*’s usefulness as a tool for comparing different denoising algorithms in order to find the one that performs the best for a specific data set.

6.1 Experimental Details

In all experiments, the DistilBert uncased model for English language (Sanh et al., 2019) provided by the HuggingFace² (Wolf et al., 2020) library is used as the prediction model. The optimization is performed with the AdamW optimizer and a learning rate of $1e-4$. We employ a cross-entropy loss accepting a probability distribution over all labels as reference input whenever appropriate (KNNTrainer, Snorkel-Trainer). Reducing this representation to a single label, i.e., standard cross-entropy, would lead to a loss of weak signals, whereas a probability allows to exploit the information from Z and T to the fullest. Each model was trained for 2 epochs (unless stated otherwise), what was enough to receive a stable result.

For the k -NN algorithm we made the experiments with different k values and provide TF-IDF features based on a dictionary of 3000 words. Hyperparameters for the WSCrossWeigh denoising algorithm are the number of folds the data will be split into, the number of partitions (that is, how many times the splitting for mistake estimation is done) and a weight-reducing rate (the value, by which the initial sample weights will be reduced to each time the sample was predicted wrongly). These parameters are tuned for each dataset individually. The following best parameter values were found empirically: 3, 10 and 0.3 for the Spam dataset, 3, 2 and 0.3 for the Spouse dataset and 2, 25, 0.7 for the IMDb dataset. Apart from that, *knodle* provides the opportunity to train the sample weights with a model different from the final classifier. In our experiments, the weights were calculated using a Bidirectional LSTM with GloVe Embeddings (Pennington et al., 2014), while the final training was performed with DistilBert using the same settings as in the experiments with other

²<https://huggingface.co/>

denoising methods. The only difference is the number of epochs on the TAC-based dataset: the best results were obtained with 1 DistilBert epoch.

6.2 Results

An overview of the numbers is given in Table 2. In the Spam dataset, all denoising methods show an improvement over the simple majority vote baseline. The data-adaptive k -NN ($k=2$) and WSCrossWeigh methods perform best in this setting. Snorkel and the standard majority voting combined with fine-tuning overfit to the noisy majority votes. This becomes obvious with the observation that Snorkel achieves a score of 0.93 with a simple logistic regression discriminative model. Interestingly, k -NN performs so well - this can serve as a proof of neighboring labels reliability.

Compared to the Spam dataset, the Spouse dataset is much larger. Since the test set is very unbalanced (the class ratio equals 0.08), all given metrics are related to the *is-spouse* relation, i.e. the underrepresented class. On average, 0.34 rules hit per instance, meaning that in almost 70% of the training data match no rule. Thus, the majority vote is mostly based on random votes, rendering a high recall but low precision. During a close inspection of the rules it became clear that they overrepresent the *is-spouse* class. Thus, the additional model training magnifies this overfitting which is expressed by an increased recall and a lower precision. The only denoising system that turned out to generalize well is Snorkel. One of the possible explanations could be that it provides an explicit rules denoising in contrast to other methods.

For IMDb, the majority vote shows that the rules have rather low quality on their own, but an additional trained model on top manages to generalize beyond the given labels. In contrast, denoising with the k -NN ($k=2$) algorithm only aggravates the problems inherent to labels of this data set, as the classifier performance drops down to a random vote (50% accuracy). This behaviour can be explained by the high density of rule hits: on average no less than 33 keywords match for each sentence, meaning that already by $k=1$, so many labels are added to its neighbors that propagation of the imprecise labelings overrules the expected benefits of k -NN. In general, there are cues that k -NN might be useful in cases where the weak labels are already rather reliable but fail in cases where weak labels are too noisy. On the contrary, WSCrossWeigh denoising algorithm, which deals not with the rules, but with the data once labeled with them, works for this dataset and can improve the results. Lastly, Snorkel

fails to calculate its discriminative model with 6800 rules in the training matrix, since its memory consumption exceeds the available limit of 32Gb RAM.

TAC-based RE dataset turned out to be the most complicated dataset among all because of both a larger size n and a larger number of rules r . Due to its specificity, there are almost no rule matches on the test set, implying that the simple majority baseline has scores close to 0. The training with DistilBert improves the results, which, however, remains considerably worse than the ones discussed above. Similar to the IMDb dataset, Snorkel denoising could not be performed on this dataset due to the amount of rules. Though in both cases we could calculate the results using the reduced number of rules, we want to avoid any data manipulations. The original k -NN algorithm also could not be performed because the computation of distances between almost 2 millions instances, which are necessary to determine the nearest neighbors, turned out to be extremely memory- and time-consuming. We work around this by applying an *approximated* k -NN algorithm instead; for example, in our experiments we used the Annoy library (Bernhardsson, 2015) and $k=3$ parameter. The poor performance of k -NN could be explained by a small average rule hits rate of the TAC-based RE data set; the possible approximation losses are also not to be neglected. In contrast, WSCrossWeigh denoising method, being rule-independent, is still able to improve the results.

7 Conclusion

This work introduces the Knowledge-Infused Deep Learning Framework *knOdle*. It provides a unified interface to work with multiple weak labeling sources, so that they can be seamlessly integrated with the training of deep neural networks. This is achieved by a tensor-based input format and an intuitive separation of weak supervision aspects and model training. Apart from that, *knOdle*'s modular approach makes it easy to add new data sets and denoising algorithms. The framework facilitates experimentation that helps researchers to derive better insights into the general correspondence between methods and weak supervision settings. From a practical perspective, *knOdle* can be used to compare different denoising methods and select the one that gives the best result for a specific task. Adding functionality to *knOdle* is straightforward. We hope that this will encourage researchers to create their own algorithms to improve learning with weakly annotated data, and incorporate them to the *knOdle* framework.

References

- Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Yangqing Jia, Rafal Jozefowicz, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dandelion Mané, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Mike Schuster, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viégas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. 2015. [TensorFlow: Large-scale machine learning on heterogeneous systems](#). Software available from tensorflow.org.
- Eugene Agichtein and Luis Gravano. 2000. [Snowball: Extracting relations from large plain-text collections](#). In *Proceedings of the Fifth ACM Conference on Digital Libraries*, DL '00, page 85–94, New York, NY, USA. Association for Computing Machinery.
- T C Alberto, J V Lochter, and T A Almeida. 2015. [TubeSpam: Comment Spam Filtering on YouTube](#). In *2015 IEEE 14th International Conference on Machine Learning and Applications (ICMLA)*, pages 138–143.
- Stephen H Bach, Bryan Dawei He, Alexander Ratner, and Christopher Ré. 2017. [Learning the Structure of Generative Models without Labeled Data](#). *CoRR*, abs/1703.0.
- E. Bernhardsson. 2015. [Annoy on github](#). Last accessed 23 April 2021.
- K. Boland and F. Krüger. 2019. Distant supervision for silver label generation of software mentions in social scientific publications. In *BIRNDL@SIGIR*.
- Oishik Chatterjee, Ganesh Ramakrishnan, and Sunita Sarawagi. 2019. [Data Programming using Continuous and Quality-Guided Labeling Functions](#). *CoRR*, abs/1911.0.
- D. Corney, M. Albakour, Miguel Martinez-Alvarez, and Samir Moussa. 2016. What do a million news articles look like? In *NewsIR@ECIR*.
- Mark Craven and Johan Kumlien. 1999. Constructing biological knowledge bases by extracting information from text sources. In *Proceedings of the Seventh International Conference on Intelligent Systems for Molecular Biology*, page 77–86. AAAI Press.
- A. P. Dawid and A. M. Skene. 1979. [Maximum likelihood estimation of observer error-rates using the em algorithm](#). *Journal of the Royal Statistical Society. Series C (Applied Statistics)*, 28(1):20–28.
- Eleazar Eskin. 2000. [Detecting errors within a corpus using anomaly detection](#). In *1st Meeting of the North American Chapter of the Association for Computational Linguistics*.
- WA Falcon. 2019. [Pytorch lightning](#). *GitHub. Note: https://github.com/PyTorchLightning/pytorch-lightning* Cited by, 3.
- Google. 2014. [Freebase data dumps](#). <https://developers.google.com/freebase/data>.
- Zhengqiu He, Wenliang Chen, Yuyi Wang, Wei Zhang, Guanchun Wang, and Min Zhang. 2020. [Improving neural relation extraction with positive and unlabeled learning](#). *Proceedings of the AAAI Conference on Artificial Intelligence*, 34:7927–7934.
- Michael A. Hedderich, Dawei Zhu, and Dietrich Klakow. 2021. [Analysing the noise model error for realistic noisy label data](#).
- Raphael Hoffmann, Congle Zhang, Xiao Ling, Luke Zettlemoyer, and Daniel S. Weld. 2011. [Knowledge-based weak supervision for information extraction of overlapping relations](#). In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pages 541–550, Portland, Oregon, USA. Association for Computational Linguistics.
- Dirk Hovy, Taylor Berg-Kirkpatrick, Ashish Vaswani, and Eduard Hovy. 2013. [Learning whom to trust with MACE](#). In *Proceedings of the 2013 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 1120–1130, Atlanta, Georgia. Association for Computational Linguistics.
- Minqing Hu and Bing Liu. 2004. [Mining and Summarizing Customer Reviews](#). In *Proceedings of the Tenth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '04, pages 168–177, New York, NY, USA. Association for Computing Machinery.
- Panos Ipeirotis, Foster Provost, and Jing Wang. 2010. [Quality management on amazon mechanical turk](#). In: *Proceedings of the ACM SIGKDD Workshop on Human Computation*.
- Sharmistha Jat, Siddhesh Khandelwal, and Partha Talukdar. 2018. [Improving distantly supervised relation extraction using word and entity based attention](#).
- Katherine A. Keith, Abram Handler, Michael Pinkham, Cara Magliozzi, Joshua McDuffie, and Brendan O'Connor. 2017. [Identifying civilians killed by police with distantly supervised entity-event extraction](#). *CoRR*, abs/1707.07086.
- Michal Kosinski, Yoram Bachrach, Gjergji Kasneci, Jurgen Van-Gael, and Thore Graepel. 2012. [Crowd iq: Measuring the intelligence of crowdsourcing platforms](#). *Proceedings of the 3rd Annual ACM Web Science Conference, WebSci'12*.
- Zornitsa Kozareva, Ellen Riloff, and Eduard Hovy. 2008. [Semantic class learning from the web with hyponym pattern linkage graphs](#). In *Proceedings of ACL-08: HLT*, pages 1048–1056, Columbus, Ohio. Association for Computational Linguistics.

- Antonios Minas Krasakis, E. Kanoulas, and G. Tsatsaronis. 2019. Semi-supervised ensemble learning with weak supervision for biomedical relationship extraction. In *AKBC*.
- Jens Lehmann, Robert Isele, Max Jakob, Anja Jentzsch, Dimitris Kontokostas, Pablo Mendes, Sebastian Hellmann, Mohamed Morsey, Patrick Van Kleef, Sören Auer, and Christian Bizer. 2014. *Dbpedia - a large-scale, multilingual knowledge base extracted from wikipedia*. *Semantic Web Journal*, 6.
- Yankai Lin, Shiqi Shen, Zhiyuan Liu, Huanbo Luan, and Maosong Sun. 2016. *Neural relation extraction with selective attention over instances*. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2124–2133, Berlin, Germany. Association for Computational Linguistics.
- Ayush Maheshwari, Oishik Chatterjee, KrishnaTeja Killamsetty, Rishabh Iyer, and Ganesh Ramakrishnan. 2020. *Data Programming using Semi-Supervision and Subset Selection*.
- Mike Mintz, Steven Bills, Rion Snow, and Daniel Jurafsky. 2009. *Distant supervision for relation extraction without labeled data*. In *Proceedings of the Joint Conference of the 47th Annual Meeting of the ACL and the 4th International Joint Conference on Natural Language Processing of the AFNLP*, pages 1003–1011, Suntec, Singapore. Association for Computational Linguistics.
- Curtis G Northcutt, Lu Jiang, and Isaac L Chuang. 2021. *Confident Learning: Estimating Uncertainty in Dataset Labels*.
- Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan, Edward Yang, Zachary DeVito, Zeming Lin, Alban Desmaison, Luca Antiga, and Adam Lerer. 2017. Automatic differentiation in pytorch. *NIPS Workshop*.
- Giorgio Patrini, Alessandro Rozza, Aditya Menon, Richard Nock, and Lizhen Qu. 2017. *Making Deep Neural Networks Robust to Label Noise: a Loss Correction Approach*.
- Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. *Glove: Global vectors for word representation*. In *Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543.
- Pengda Qin, Weiran Xu, and William Yang Wang. 2018. *Robust distant supervision relation extraction via deep reinforcement learning*. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2137–2147, Melbourne, Australia. Association for Computational Linguistics.
- Alexander Ratner, Stephen H Bach, Henry R Ehrenberg, Jason Alan Fries, Sen Wu, and Christopher R é. 2017. *Snorkel: Rapid Training Data Creation with Weak Supervision*. *CoRR*, abs/1711.1.
- Vikas C. Raykar and Shipeng Yu. 2012. *Eliminating spammers and ranking annotators for crowdsourced labeling tasks*. *Journal of Machine Learning Research*, 13(16):491–518.
- Ines Rehbein and Josef Ruppenhofer. 2017. *Detecting annotation noise in automatically labelled data*. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1160–1170, Vancouver, Canada. Association for Computational Linguistics.
- Sebastian Riedel, Limin Yao, and Andrew McCallum. 2010. Modeling relations and their mentions without labeled text. In *Proceedings of the 2010 European Conference on Machine Learning and Knowledge Discovery in Databases: Part III, ECML PKDD’10*, page 148–163, Berlin, Heidelberg. Springer-Verlag.
- Sebastian Riedel, Limin Yao, Andrew McCallum, and Benjamin M. Marlin. 2013. *Relation extraction with matrix factorization and universal schemas*. In *Proceedings of the 2013 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 74–84, Atlanta, Georgia. Association for Computational Linguistics.
- Alan Ritter, Luke Zettlemoyer, Mausam, and Oren Etzioni. 2013. *Modeling missing data in distant supervision for information extraction*. *Transactions of the Association for Computational Linguistics*, 1:367–378.
- Brendan van Rooyen, Aditya Krishna Menon, and Robert C Williamson. 2015. *Learning with Symmetric Label Noise: The Importance of Being Unhinged*.
- Benjamin Roth. 2014. *Effective distant supervision for end-to-end knowledge base population systems*. Ph.D. thesis, Saarland University.
- Benjamin Roth and Dietrich Klakow. 2013. *Feature-based models for improving the quality of noisy training data for relation extraction*. In *Proceedings of the 22nd ACM International Conference on Information and Knowledge Management, CIKM ’13*, page 1181–1184, New York, NY, USA. Association for Computing Machinery.
- Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2019. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *ArXiv*, abs/1910.01108.
- Yuming Shang. 2019. *Are noisy sentences useless for distant supervised relation extraction?*
- Kai Shu, Subhabrata Mukherjee, Guoqing Zheng, Ahmed Hassan, Milad Shokouhi, and Susan Dumais. 2020. *Learning with weak supervision for email intent detection*. pages 1051–1060.
- Snorkel. 2020a. *Detecting spouse mentions in sentences*. Last accessed 25 February 2021.
- Snorkel. 2020b. *Snorkel intro tutorial: Data labeling*. Last accessed 25 February 2021.

- Rion Snow, Daniel Jurafsky, and Andrew Ng. 2004. Learning syntactic patterns for automatic hypernym discovery. In *NIPS*, volume 17.
- Lucas Sterckx, Thomas Demeester, Johannes Deleu, and Chris Develder. 2014. Using active learning and semantic clustering for noise reduction in distant supervision. In *NIPS 2014*.
- Sainbayar Sukhbaatar, Joan Bruna, Manohar Paluri, Lubomir Bourdev, and Rob Fergus. 2015. [Training Convolutional Networks with Noisy Labels](#).
- Ang Sun and Ralph Grishman. 2012. [Active learning for relation type extension with local and global data views](#). In *Proceedings of the 21st ACM International Conference on Information and Knowledge Management*, CIKM '12, page 1105–1112, New York, NY, USA. Association for Computing Machinery.
- M. Surdeanu. 2013. Overview of the tac2013 knowledge base population evaluation: English slot filling and temporal slot filling. *Theory and Applications of Categories*.
- Shingo Takamatsu, Issei Sato, and Hiroshi Nakagawa. 2012. [Reducing wrong labels in distant supervision for relation extraction](#). In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 721–729, Jeju Island, Korea. Association for Computational Linguistics.
- Philippe Thomas, Illés Solt, Roman Klinger, and Ulf Leser. 2011. [Learning protein–protein interaction extraction using distant supervision](#). In *Proceedings of Workshop on Robust Unsupervised and Semisupervised Methods in Natural Language Processing*, pages 25–32, Hissar, Bulgaria. Association for Computational Linguistics.
- Stephen Tratz and Eduard Hovy. 2010. A taxonomy, dataset, and classifier for automatic noun compound interpretation. In *Proceedings of the 48th Annual Meeting of the Association for Computational Linguistics*, pages 678–687.
- Paroma Varma, Frederic Sala, Ann He, Alexander Ratner, and Christopher Ré. 2019. [Learning Dependency Structures for Weak Supervision Models](#).
- Zihan Wang, Jingbo Shang, Liyuan Liu, Lihao Lu, Jiacheng Liu, and Jiawei Han. 2019. [CrossWeigh: Training named entity tagger from imperfect annotations](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5154–5163, Hong Kong, China. Association for Computational Linguistics.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander M Rush. 2020. [HuggingFace’s Transformers: State-of-the-art Natural Language Processing](#).
- Fei Wu and Daniel S. Weld. 2007. [Autonomously semantifying wikipedia](#). In *CIKM*, page 41–50. ACM.
- Wei Xu, Raphael Hoffmann, Le Zhao, and Ralph Grishman. 2013. [Filling knowledge base gaps for distant supervision of relation extraction](#). In *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 665–670, Sofia, Bulgaria. Association for Computational Linguistics.
- Limin Yao, Aria Haghighi, Sebastian Riedel, and Andrew McCallum. 2011. [Structured relation discovery using generative models](#). In *Proceedings of the 2011 Conference on Empirical Methods in Natural Language Processing*, pages 1456–1466, Edinburgh, Scotland, UK. Association for Computational Linguistics.
- Yuhao Zhang, Victor Zhong, Danqi Chen, Gabor Angeli, and Christopher D. Manning. 2017. [Position-aware attention and supervised data improve slot filling](#). In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing (EMNLP 2017)*, pages 35–45.
- Zhilu Zhang and Mert R Sabuncu. 2018. [Generalized Cross Entropy Loss for Training Deep Neural Networks with Noisy Labels](#).